

Sentiment Analysis of Tweets to Predict Political Affiliation

Abstract

This paper delves into code written to find the political polarity of tweets through different methodologies. Specifically, it is able to find political polarity by pulling tweets directly from Twitter. To initially analyze these different methodologies, a training data set of tweets was found that contained tweets and the associated polarity of those writing the tweets. This data set was used to analyze the connection between the tweet and the associated political party of the source. After the tweets and associated party were extracted, different models were used to measure political polarity, including linear SVM classifiers, naive bayes, logistic regression, and LSTM. The results of all of these pipelines were analyzed and it was found that linear SVC had the highest accuracy of 0.85. This highest performing methodology was then used to create a pipeline to predict the polarity of tweets coming from Twitter itself. The political parties that were used for classification included Republican (R) Democrat (D) or Independent (I).

Introduction:

When reading a tweet online or glancing through a news article, the political polarity of the type of writing may be unclear. Knowing this information could optimize a user's experience by helping ensure that a person receives quality information and help make the person more aware of where the information they are consuming is coming from. Furthermore, knowing the political polarity of different texts read online can help create a more pleasurable reading experience for the user. Knowing this information can help them decide if they want to continue reading the tweet/article or not. TF-IDF was implemented in conjunction with linear SVM classifiers to eventually take in a tweet as input and output the political party that the user is affiliated with.

State of the art:

Currently, there are different studies that have been made to help reach a similar goal of finding the political polarity of different tweets. Although there is no mainstream feature that finds this, there are different studies that have done research on this topic. Many of these studies relate to finding political polarity of a specific subset of users or focus on a specific time period. This paper focuses on the methods explored to create a generalized political polarity of any tweet given a hashtag. It also explores this feature using different methodologies and analyzes some of the highlights of each one. Specifically, the methods used to find polarity were TF-IDF along with linear SVC classifiers, LSTM and naive bayes

Solution::

Before the code for the prediction pipeline was created, the appropriate training data set was needed. Finding this was done through searching through various data sets of tweets, and ensuring that the one selected included the political party of the source. Searching for this attribute of the dataset was necessary to ensure that there was a way to verify the results of the pipeline. Once the data set of senators was found, which matched these requirements, the data set was visualized to get further information on the data set. Specifically, the division of those in both

parties was visualized into a bar graph, and it was found that the majority of those in the data set identified as Republican, and that the party with the lowest number of senators was Independent (Figure 1). The data set was further visualized to show the number of tweets from senators from each state (Figure 2). It was found that senators from Alabama tweeted the least, at around 2000 tweets, and those from more active states had over 6000 tweets.

Once the data was analyzed, to analyze the different models, the raw data text was needed. Therefore, the data set was split so that the actual tweets were contained in the variable X and the party for corresponding tweets was contained in variable Y. Afterwards, the training data set was split into 4 parts so that each model had a different subset of the data to work on. Then, to be able to use the tweets and corresponding data in the models selected, the TfidfVectorizer function was used and stop words were removed. This function converted the strings from the data set into vectors that could be used within the models.

The first model that was used to find the political polarity of the data was naïveté bayes. This model was applied to the split data set and the results of the model were analyzed. Then a similar process was performed on the model called linear SVC, using part of the training data and the results were analyzed. This process was repeated again for logistic regression, and results were analyzed.

Results:

After training the model for several hours, we were able to achieve a 85% test set accuracy , essentially surpassing our initial target goal of 80%.

Table 1: Examples of Correctly Classified Tweets

Tweet	Prediction
How could the Obama administration pay Iran \$1.7 billion out of a fund that should already be depleted?	Republican
Obamacare is giving us higher premiums and canceled policies -- much bigger problem than the website. trainwreck	Republican
SCOTUS should strike down TX law that makes it much more difficult for women to make their own health care choices.	Democrat
Cyber is a national security imperative, essential to protecting Maine infrastructure—like our electric grid and healthcare system—from bad actors.	Independent

Table 2: Examples of Misclassified Tweets

Tweet	Prediction	Actual
Flynn's resignation is tip of iceberg. So many Qs unanswered about Russia's interference in election and its communication with Trump team	Democrat	Republican

Proud to join with Senator Shaheen at York Hospital to call for funding to fight opiate epidemic	Democrat	Independent
My bill to help recruit and retain federal workers would support law enforcement, permitting, agriculture and many other jobs and biz across ND	Republican	Democrat
Jim Marshall bravely served during WWII. It was an honor to deliver the medals he earned to his family yesterday.	Republican	Democrat

Related work:

Certain papers have also found political polarity of tweets for a specific subset of users through related mechanisms. A specific study done found the political inclination of users during the 2016 election, and whether they were more likely to support Trump or Hillary. It also then went further to measure political homophily after classifying the Twitter users.

The first part of the study analyzed a data set of 4.9 million tweets of 18450 users. These tweets were classified based on their likelihood of voting for either of the two nominees, and the categories that reflected these political polarities were named as: whatever, Trump supporter, Hillary supporter, positive, neutral and negative. One way that the study determined if the tweet was considered political and was researched more, is that certain indicators were looked for. Specifically, if the tweet was a candidate's retweet or it targets at least one candidate, that was used to label the tweet as "political". Additionally, if the tweet mentioned at least one candidate or contained a candidate's name, it was also considered a political tweet. After this sentiment analysis was performed on the selected tweets. It was then further found that if a tweet mentioned a specific candidate, then that tweet was affiliated with the candidate mentioned. However, for the tweets that mentioned both of the candidates, two different strategies were implemented to find the user's political affiliation, and the Stanford Parser tool was used to determine how the user felt about each candidate in this situation.

Additionally, to further find the political polarity of users, the SentiStrength tool was used, which searches for words that it already "knows" to be positive or negative. After searching and finding these words in the sentence, values that indicate how positive or negative the word is are marked, and the difference between these two values indicates how positive or negative the final tweet is. The average sentiment for all of the users was found. For the values that were found, the average sentiment that fell into a specific range was marked into one of the 6 categories as mentioned above.

The approaches found in the study mentioned are similar to the approaches of sentiment analysis as mentioned in this paper. Specifically, the SentiStrength tool assigns a value to each word indicating positive or negative sentiment, and the code discussed in this paper acts similarly.

A different study that is similar to the one mentioned in this paper is a study that focuses on tweets related to the Indian general election of 2019 and studies sentiments of Twitter users towards the major political parties in India. 3869 total tweets were analyzed and then classified into two parties, and any tweets that did not fall into these two parties were categorized into "other".

Conclusion

After carefully preprocessing, debugging, training, and fine-tuning the model, Linear SVM ended up with the highest testing accuracy of 85%, which has passed the target goal of 80% accuracy. After downloading and analyzing the dataset, there was a concern for its vastness and variability compared to datasets by related studies. As such, the results were not expected to be as comparable to previous work, whose top-performing model had 91.6% accuracy. It is believed that this accuracy, as well as the final accuracy, is far higher than most average humans' accuracy in predicting the political party of a random U.S senator, given one random tweet without context.

While it is a thrill that the target goal was met, if given more time, the results could have been even closer to the stretch goal of 90% accuracy. The main barrier was certainly the extensive amount of time that it took to run the model in its entirety, as it is apparent that further hyperparameter fine-tuning could have improved the accuracy. Another step that could have been taken, if there was more time, is to further investigate/analyze the dataset. With a bit more manual investigation, There would have been swift removal of tweets whose data seemed to be out of the ordinary (for instance, tweets that contained emoticons).

The main challenges throughout this project were determining how to best preprocess the data and constructing the model in Tensorflow, which was a framework new to all of us. The US senators' Tweets dataset was also very large and took many hours to train and output models from. There was an initial inquiry to include mentions in the vocabulary list, according to the hypothesis that including certain people would provide relevant information as to the political beliefs of the tweeter; with all of this data, the vocab size was around 286,000. No one in the group had much experience with Keras prior to implementing this model, so learning how to design, create, and train models in Keras and Tensorflow proved to be a steep learning curve for us. Familiarizing with the Keras framework was quite time consuming, and even after learning how to transfer the TensorFlow model into Keras, debugging in this unfamiliar framework was challenging. The first struggle came with Tensorflow, as the input data were improper types for many of the layers/functions. Simply getting accustomed to the differences between TensorFlow and Keras was a very consuming and challenging step in developing the model. because of the sheer amount of data that the algorithm was attempting to train the model on, it took a very long time to fine-tune hyperparameters, even when training on SGD. However, there was still room to test various hyperparameters, including learning rate, batch size, more dense layers, momentum, and embedding size. After many iterations, the code was settled on a learning rate of 0.1, a batch size of 64, momentum of 0.8, and an embedding size of 32.

In addition, there was an original interest within the group in seeing how the model would be used to create a real-time tweet aggregator that can fetch tweets through an API and, targeting a specific user or keyword, identify what political party that tweet is likely affiliated with and filter those tweets in a live feed. There was also some dabbling with the idea of doing news articles, not tweets, to create some real-time new aggregator for a specified news site.

There was no anticipation that acquiring the data and running the model would take as many days/hours as they did, so there unfortunately was not enough time to explore this task. If the opportunity for a similar project came again, it would definitely be interesting to attempt a real-time tweet aggregator. Other than not reaching this last step, the program stuck pretty closely to the original plan for implementing this project, and while many parts took longer than expected, the group is overall pleased with the trajectory of the project.

The group learned a multitude of new skills throughout this project. From interacting with massive amounts of data, to learning the Keras framework, to implementing a Tensorflow model based on a scientific paper, there was lots of exposure to many new experiences throughout this process. The group also amassed much practice in debugging a model on the fine-tuning parameters to achieve the target results without any TA assistance, simulating what it would be like to design a Machine Learning algorithm outside of this course. This project not only helped appreciate the power of Machine Learning in general, but specifically showed the potential of NLP models in the realm of social media. This project goes to show how much coded language is used on social media to further party agendas. The fact that it can become so easy for a deep learning network to determine posters' political party affiliation from a single tweet about nearly any topic demonstrates how much political content, and frankly, propaganda, people consume on a daily basis. It truly emphasizes the importance of being vigilant users of social media.

References

- Sitaram Asur and Bernardo A. Huberman. 2010. Predicting the future with social media. Computing Research Repository.
- Gustavo Laboreiro, Lu'is Sarmiento, Jorge Teixeira, and Eug'enio Oliveira. 2010. Tokenizing microblogging messages using a text classification approach. In AND '10: Proceedings of The Fourth Workshop on Analytics for Noisy Unstructured Text Data, New York, NY, USA, October. ACM.
- Wang H, Can D, Kazemzadeh A, Bar F, Narayanan S
A system for real-time Twitter sentiment analysis of 2012 US Presidential election cycle., ACL (System Demonstrations) (2012), pp. 115-120
- Conover M, Gonçalves B, Ratkiewicz J, Flammini A, Menczer F. (2011a) "Predicting the political alignment of Twitter users." In Proceedings of SocialCom/PASSAT Conference, pp. 192–199
- Conover M, Ratkiewicz, J, Francisco M, Goncalves B, Menczer F, Flammini A. (2011b) "Political Polarization on Twitter." In International AAAI Conference on Web and Social Media.
- Choy M, Cheong MLF, Laik MN, Shung KP (2011) "A sentiment analysis of Singapore presidential election 2011 using twitter data with Census correction." CoRR abs/1108.5520
- Mishra P, Rajnish R, Pankaj Kumar
Sentiment Analysis of Twitter Data: Case Study on Digital India. INCITE-2016, Amity University, India (2016)
- Marchetti-Bowick M, Chambers N. (2012) "Learning for microblogs with distant supervision: political forecasting with Twitter." In Proceedings of the EACL Conference, pp. 603–612
- K.Kohnson, D. Jin, D. Goldwasser. *Leveraging Behavioral and Social Information for Weakly Supervised Collective Classification of Political Discourse on Twitter*. ACL, 2017. Github repository, <https://github.com/kmjohanson/twitter-framing>
- E. Summers, N. Ruest. *Twarc*. University of Maryland, 2017. GitHub repository, <https://github.com/DocNow/twarc>
- C. Olah. *Understanding LSTM Networks*. 2015. Github repository <http://colah.github.io/posts/2015-08-Understanding-LSTMs/>.